



H P C

High Performance Computing

Agenda

- **What do we do**
 - Definitions, terms, technologies, motivations & examples...
- **Computer Vision:**
 - Classic algorithm
 - Deep Learning

What Do We Do?

HPC Definition

 The use of parallel processing for running advanced application programs efficiently and quickly:

- *Speed* - reducing time to solution
- *Energy efficiency* - doing more with less power
- *Upscaling* - handling larger problems
- *High throughput* - handle large volumes of data in real-time

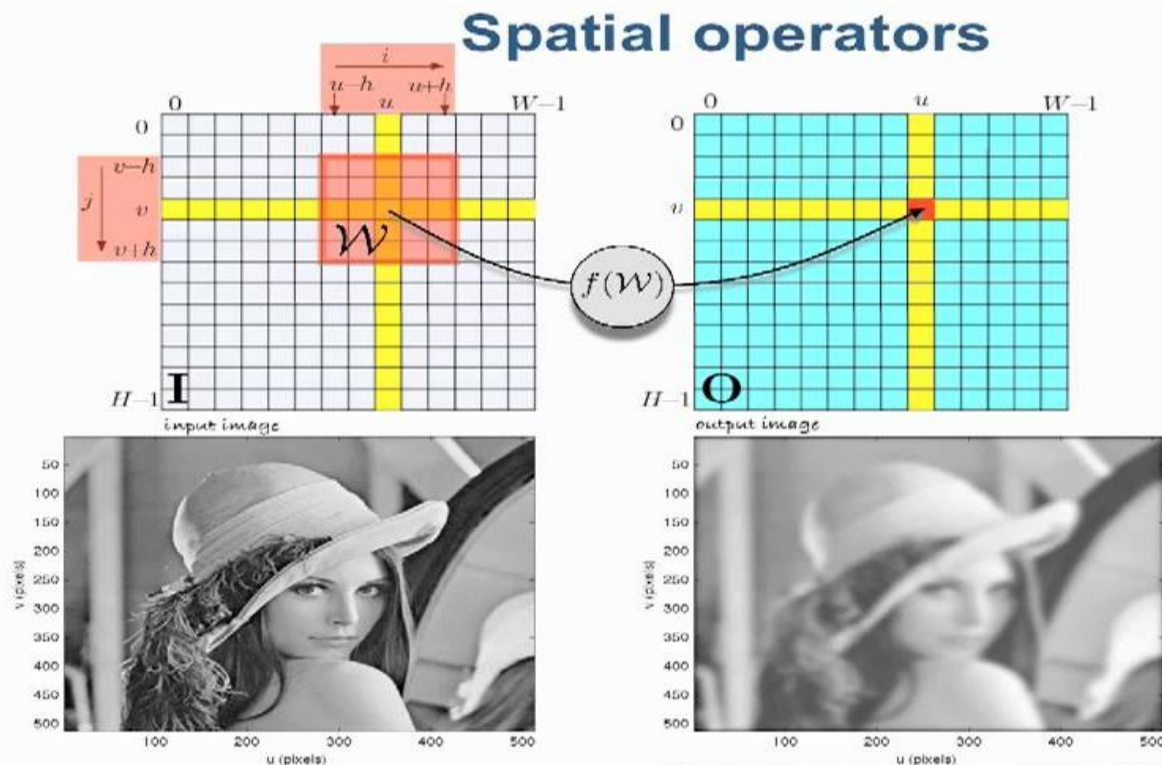
What Do We Do?

Use HPC capabilities for Image Processing

Large set of data (Large frame = Lots of pixels)

Same operation \ calculation on each datum (pixel)

=> Parallel computation is self-evident



What Do We Do?

In numbers:

🚩 HD frame = $1920 \times 1080 \approx 2\text{MPixel}$

3 color components: RGB

=> 6MB

🚩 SkEye Sensor = $10,000 \times 7096 \approx 70\text{MPixel}$

2 bytes per component

=> **140 MB**

🚩 Matlab + CPU processing

~ 2 seconds

🚩 CUDA/OpenCL + GPU processing

~ several milliseconds



What Do We Do?

Our Main Tool Set

🚧 CPU & GPU – Hardware

🚧 OpenCL – Development framework over any GPU

🚧 CUDA – Development over NVIDIA GPUs


🚧 Third party

- Boost C++

- OpenCV



OpenCL



GPU - Graphics Processing Unit

- Originally designed for graphic implementations
- Contains large amount of cores, designed to work in parallel
- GPGPU = General Purpose GPU



OpenCL - Open Computing Language

- **Open standard language for parallel computation on heterogeneous platforms (CPU, GPU, FPGA, etc.)**
- Producing small programs / calculations (Kernel) that work concurrently on “work items”
- Useful for working on all pixels of an image in parallel (each one considered a work item).
- Features
 - Dynamic library implemented by the vendor
 - Syntax resembling C language with APIs to control the GPU


K H R  N O S[®]
G R O U P



OpenCL

CUDA - Compute Unified Device Architecture

- **Parallel computation on NVIDIA GPUs (non standard).**
- **Provides better fit and higher acceleration on NVIDIA GPUs.**
- **Tools and libraries provided by NVIDIA**
 - Common algorithms
 - FFT, convolutions, and other complicated methods
 - Use of tensor cores (to support deep learning)
 - Multi GPU support
- **NVIDIA GPUs provide limited support in OpenCL, some platforms don't.**



TESLA

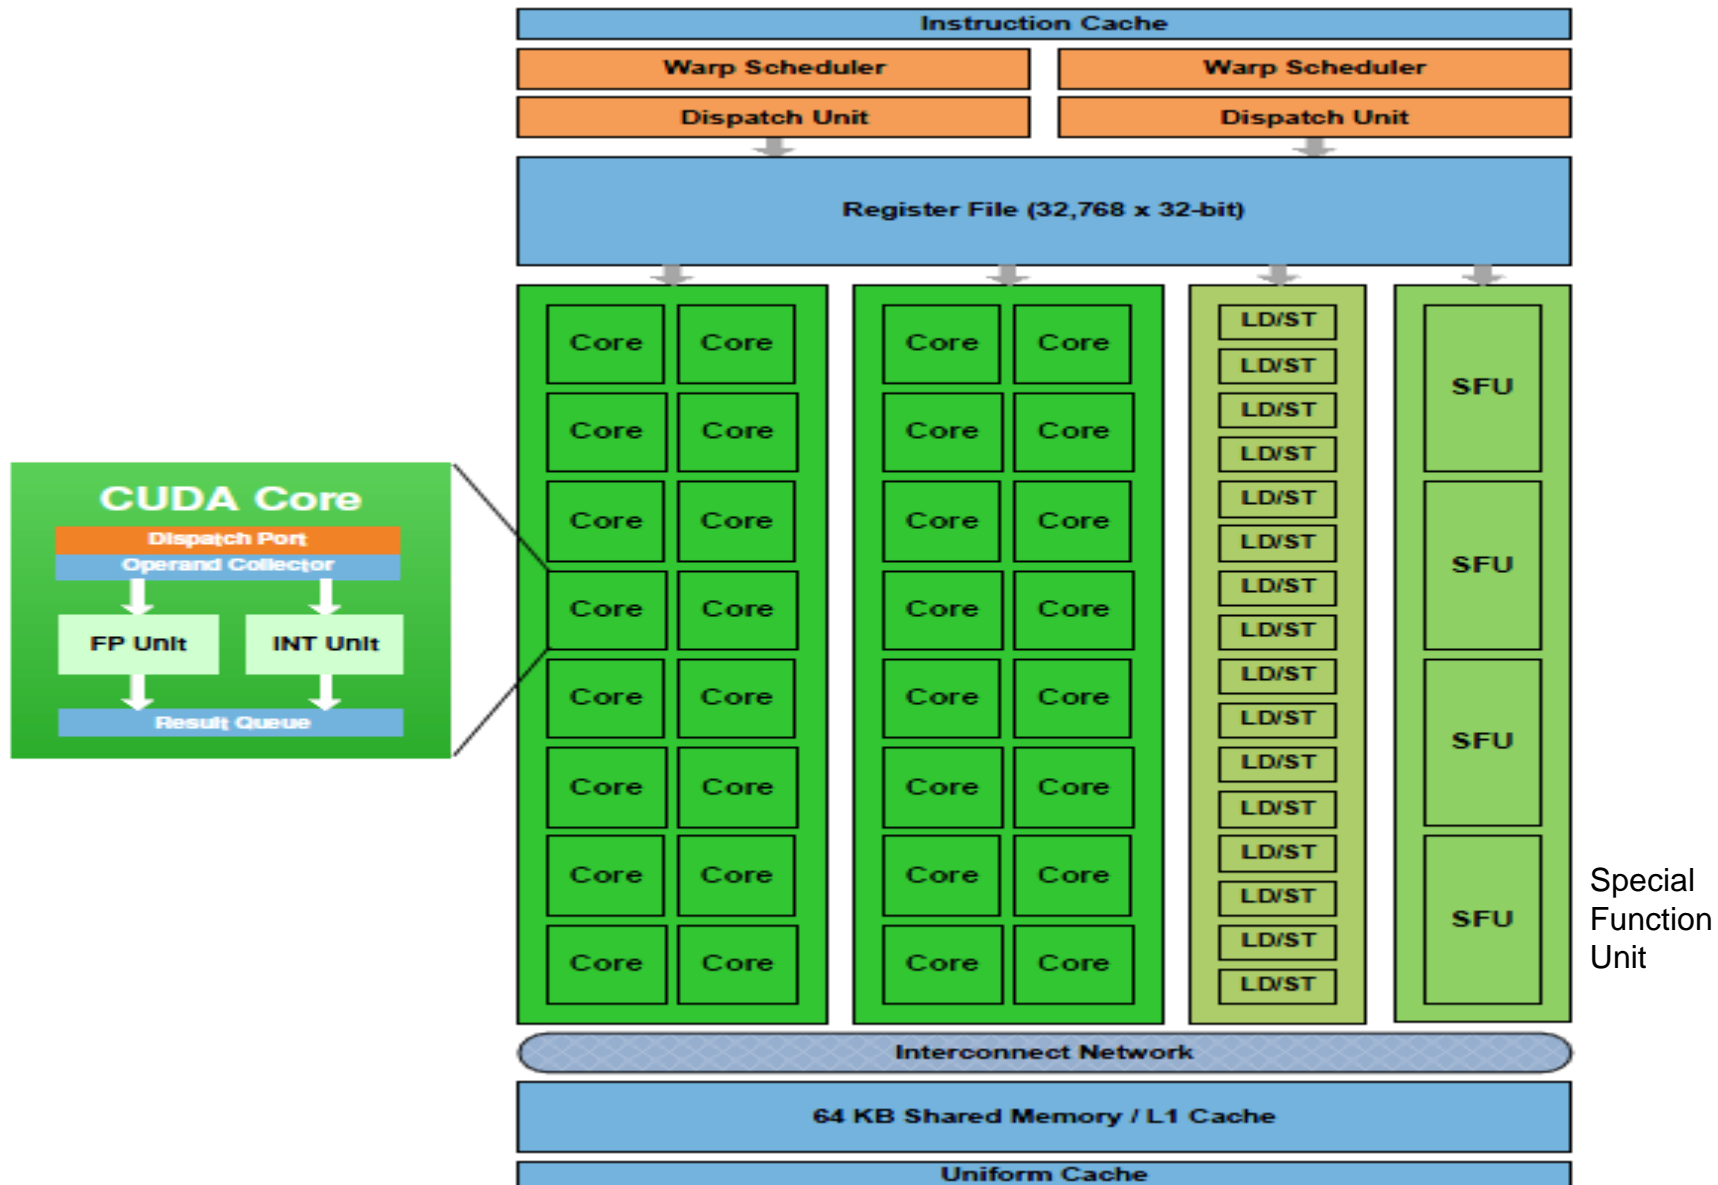


QUADRO



GEFORCE

CUDA – Micro architecture



Parallel computation example

Adding two vectors



Add vector of N elements

Assume global work size of
Nx1

CPU:

```
void vecAdd(const int N,  
const float * a,  
const float * b,  
float * c)  
{  
    for (int i=0; i< N; ++i)  
    {  
        c[i] = a[i] + b[i];  
    }  
}
```

GPU (openCL):

```
__kernel void vecAdd(  
    __global const float * a,  
    __global const float * b,  
    __global float * c)  
{  
    int i = get_global_id(0);  
    c[i] = a[i] + b[i];  
}
```

Parallel comp example: Array Sum

Sum of all elements of array

array elements




$$\begin{aligned}\text{SUM} &= a[0] + a[1] + a[2] + a[3] + a[4] \\ &= 5 + 2 + 7 + 9 + 6 \\ &= 29\end{aligned}$$

© w3resource.com

CPU



GPU



Two projects we've done

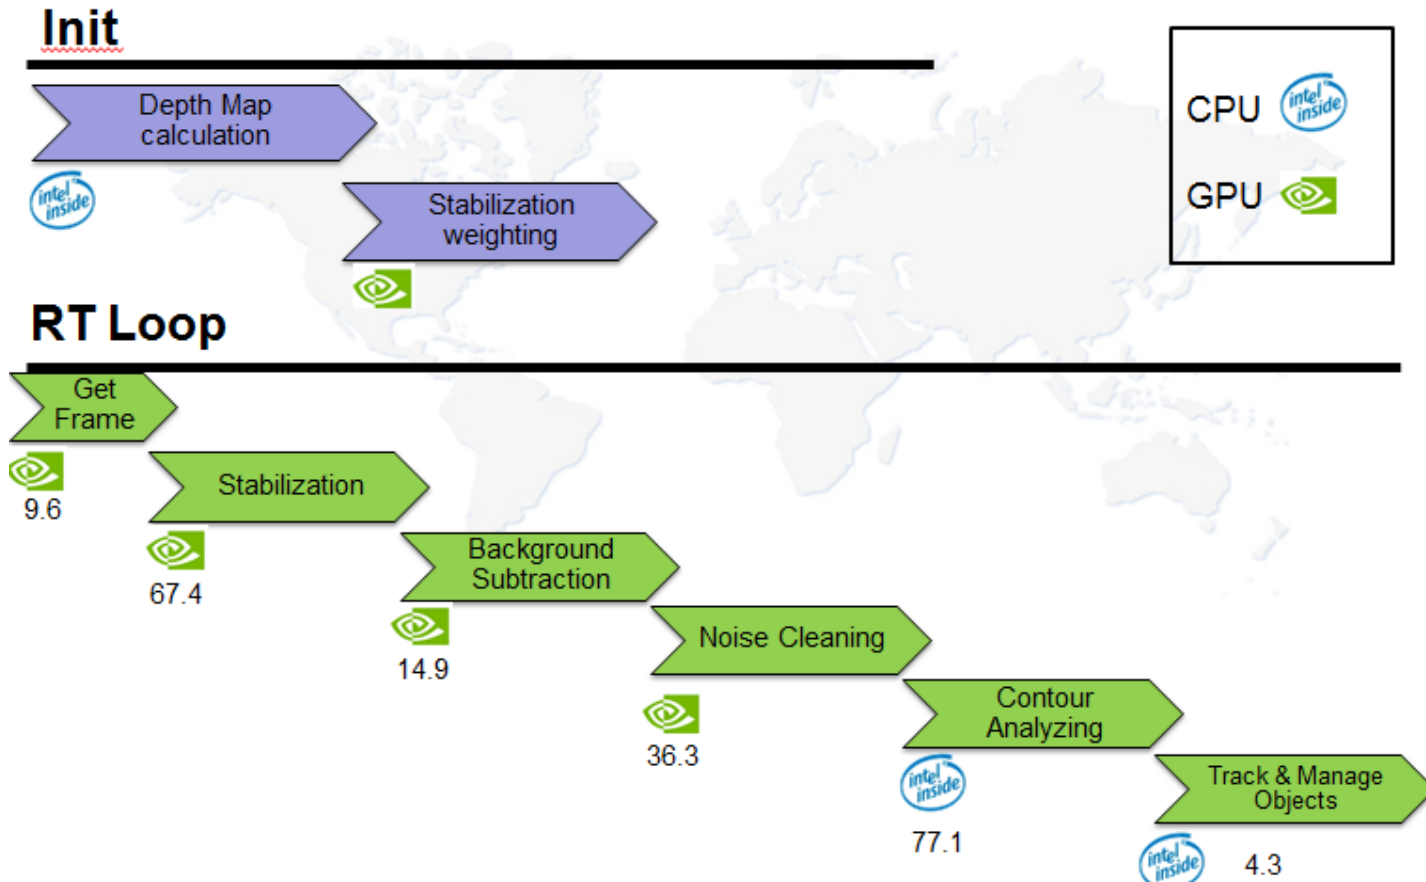
GroundEye – Classic algorithm

DLEWARE – Deep Learning

GroundEye: Video Motion Detection

▶ VMD: Video Motion Detection for 2 types of objects at various distances

- ▶ An intensive image stabilizer corrects movements of the camera.
- ▶ By sampling the background we detect the blobs that are different
- ▶ Objects movement is analyzed to filter out jittering objects (e.g. trees)



GroundEye: VMD stabilization

Get
Frame

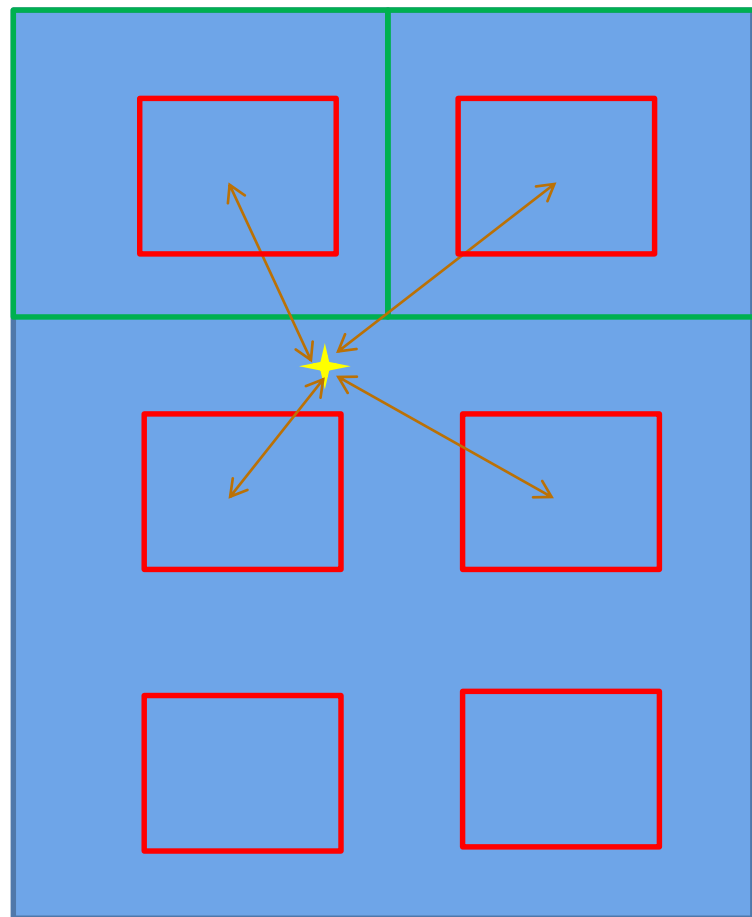
Stabilization

Background
Subtraction


Noise
Cleaning

Contour
Analyzing

Track & Manage
Objects



GroundEye: VMD stabilization



Get
Frame

Stabilization

Background
Subtraction

Noise
Cleaning

Contour
Analyzing

Track & Manage
Objects



GroundEye: VMD object detection

Get
Frame



Stabilization

Background
Subtraction


Noise
Cleaning

Contour
Analyzing

Track & Manage
Objects



GroundEye: VMD objects tracking



Get
Frame

Stabilization

Background
Subtraction

Noise
Cleaning

Contour
Analyzing

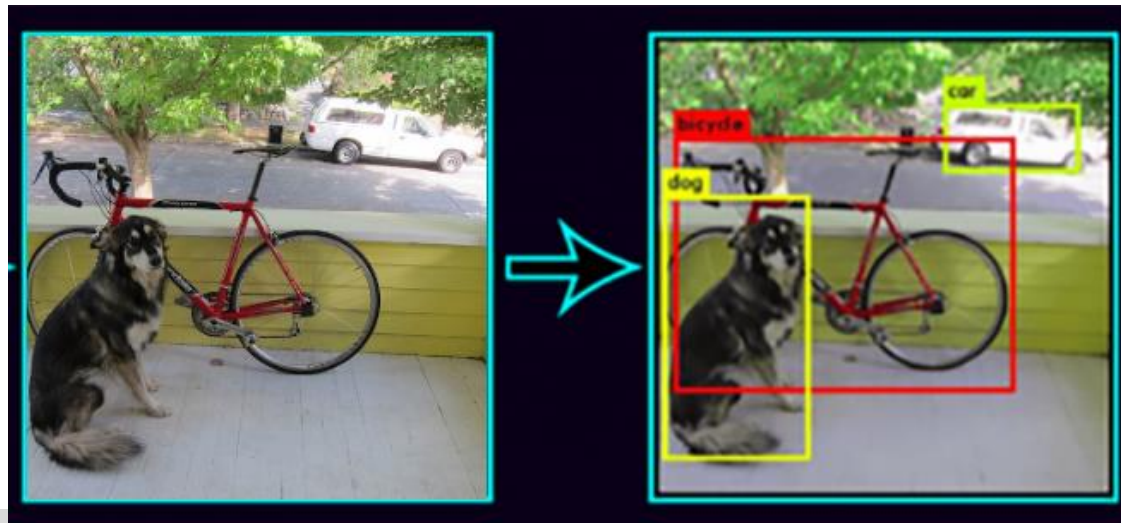
Track & Manage
Objects



DLEWARE: Deep Learning Software

**Given an input video frame,
detect all instances from the classes it had been trained to
detect,
delivering their characteristics:**

- 📶 Identification confidence
- 📶 Type confidence
- 📶 Size: Width & Height
- 📶 Position: CenterX & CenterY



Method

🚩 Unlike classic algorithms,

that dictate what each one of the image categories of interest look like directly in code: Mature, proven, and optimized, but not versatile.

🚩 In Deep Learning,

you provide the computer with many examples of each image class and develop a way to deduce from them the visual appearance of each class. can offer greater accuracy and versatility but demands large amounts of computing resources.



🚩 To do that,

- first accumulate a training dataset of labeled images (**Training phase**),
- then feed it to the computer to process the data (**Inference Phase**).

Inference Phase is where HPC comes into the picture!

HW & SW resources

Three platforms:


	PC#1	PC#2	<u>Embedded Jetson</u>
Operating system	Windows 	Linux (Ubuntu) 	L4T – Linux For Tegra
CPU	<u>Intel Xeon</u>	Intel Xeon	ARM
GPU	<u>NVIDIA GeForce GTX 1080</u>	NVIDIA GeForce GTX 1080	NVIDIA Tegra
DL framework	<u>Tensorflow</u> (Built from sources)	Tensorflow (TBC)	Tensorflow (TBC)
DL Network	<u>CNN – YoloV3</u>	CNN – YoloV3	CNN – YoloV3
Programming Language	<u>C++ & CUDA</u>	C++ & CUDA	C++ & CUDA
Third party SW	<u>TensorRT</u> , CuDNN, <u>Boost</u> , <u>OpenCV...</u>	TensorRT, CuDNN, Boost, OpenCV...	TensorRT, CuDNN, Boost, OpenCV...



Thank You!

Q&A

DLEWARE – Deep Learning Engine



Capture Image



Adapt image size & format



YOLOv3 on
Tensorflow or
TensorRT




Translate
tensor output



Filter out
low confidence
or high overlap



Undo
adaptations
and send result



DLEWARE – Deep Learning Engine



Thank You!

Q&A